

Report on Cut-in Vehicle Detection

Date of Submission: 13/07/2024

Name: Aagam Chhajer

Institute: SRM Institute of Science and Technology Kattankulathur

Branch/Specialization: B.TECH AI

Registration Number: RA2211047010110

Internal Mentor: Dr. Sumathy G

External Mentor: Dr. Vasudha Kumari (AI Software Solutions Engineer, Intel)

Problem Statement.

Autonomous driving systems are becoming increasingly common, and a crucial aspect of their development involves ensuring the safety and reliability of vehicle detection and predicting their behavior. One particularly challenging scenario in autonomous driving is when vehicles, such as cars, motorcycles, bicycles, and even pedestrians, abruptly cut into the path of the autonomous vehicle. Detecting these sudden intrusions accurately and quickly is essential to prevent potential accidents and ensure smooth navigation.

The goal of this project is to build a machine learning model using the YOLO framework to detect sudden vehicle cut-ins. The model will be trained on the India Driving Dataset (IDD) Temporal, which provides a wide range of driving scenarios typically encountered in Indian traffic conditions.

Technical Approach.

To tackle the task of detecting sudden vehicle cut-ins in autonomous driving scenarios, we utilized the YOLO model, renowned for its advanced capabilities in real-time object detection. Our approach was built upon the India Driving Dataset (IDD) Temporal, which provides a comprehensive range of typical driving scenarios observed in Indian traffic. This dataset includes diverse instances featuring various types of vehicles and pedestrians, crucial for training our model to effectively recognize and respond to sudden intrusions on the road.

Data Preparation.

Initially, we focused on preparing our dataset to adhere to the YOLO format requirements. This involved organizing the dataset into distinct directories for images and their respective labels, segregated for both training and validation purposes. Our dataset structure was set up as follows:

- Training images: /dataset/images/train
- Validation images: /dataset/images/val
- Training labels: /dataset/labels/train

- **Validation labels:** `/dataset/labels/val`

Each label file within these directories contained essential metadata such as bounding box coordinates and class identifiers corresponding to the objects depicted in the images.

Subsequently, we crafted a configuration file named `idd_temporal.yaml` to formalize the dataset's organization and define crucial parameters. This configuration file specifically outlined:

- Paths leading to the training and validation datasets (`/dataset/images/train`, `/dataset/images/val`, `/dataset/labels/train`, `/dataset/labels/val`).
- The total number of distinct classes present within our dataset, encompassing categories such as cars, motorcycles, and pedestrians.
- Clear and concise names assigned to each class, ensuring accurate identification during model training.

This meticulous approach ensured that our dataset was meticulously structured and perfectly poised for training our YOLO model, establishing a robust framework for subsequent stages of development and optimization.

Model Training.

During the training phase, we capitalized on the transfer learning capabilities of the YOLO model by initializing it with pre-trained weights. This approach significantly expedited the training process and bolstered performance, leveraging the model's prior training on a diverse and extensive dataset.

Our training script was structured to orchestrate the training loop, specifying key parameters such as:

- **Epochs:** We set the training to run for 50 epochs, allowing the model to iteratively learn from the dataset.
- **Batch Size:** Each iteration processed a batch of 16 samples, balancing computational efficiency with gradient accuracy.
- **Data Augmentation:** Techniques like random scaling and cropping were employed to augment the dataset, enhancing the model's ability to generalize across various scenarios.
- **Optimization:** Stochastic gradient descent (SGD) was used to optimize the model's parameters, fine-tuning its performance with each epoch.

Throughout the training process, we meticulously monitored the model's performance using the validation dataset. Key metrics such as loss, precision, recall, and mean Average Precision (mAP) were regularly evaluated. This iterative evaluation allowed us to gauge the model's efficacy in detecting sudden vehicle cut-ins and make informed adjustments to the training strategy as needed, ensuring continual improvement in performance and accuracy.

Model Evaluation.

After completing the training phase, our next step was to evaluate the model's performance using the validation dataset to assess its accuracy in detecting sudden

vehicle cut-ins and other objects within driving scenes. The evaluation script was designed to compute several crucial metrics:

- **Precision:** This metric measured the proportion of correctly predicted positive instances (true positives) out of all instances predicted as positive, indicating the model's ability to minimize false positives.
- **Recall:** Recall quantified the model's capability to correctly identify all positive instances (true positives) out of all actual positive instances, illustrating its sensitivity to detecting true positives.
- **mAP@0.5:** Mean Average Precision at an IoU (Intersection over Union) threshold of 0.5 provided an average measure of precision across all classes, indicating the model's performance in object localization.
- **mAP@0.5:0.95:** Mean Average Precision across multiple IoU thresholds (from 0.5 to 0.95) offered a broader evaluation of the model's accuracy in object detection across various levels of overlap between predicted and ground truth bounding boxes.

These metrics collectively offered a comprehensive assessment of the model's effectiveness in detecting different classes of vehicles and pedestrians. High precision indicated minimal false positives, while high recall demonstrated the model's proficiency in capturing true positives. The mAP scores validated the model's robust performance under diverse conditions, showcasing its reliability in real-world applications.

Following the evaluation on the validation dataset, we proceeded to validate the real-world applicability of our trained YOLO model by conducting testing on a separate set of test images. Here's an overview of the testing process:

1. **Testing Script:** We developed a testing script specifically tailored for running inference on the test images. This script utilized the trained YOLO model to predict bounding boxes and assign class labels to detected objects within each image.
2. **Inference Process:** During testing, the model processed each test image, applying its learned detection capabilities to identify and localize objects such as vehicles and pedestrians. The output included bounding box coordinates and corresponding class labels for each detected object instance.
3. **Visualization:** To facilitate clear interpretation and assessment of the model's performance, we employed tools like OpenCV and Matplotlib for visualizing the results. These libraries allowed us to overlay bounding boxes on the test images, providing a visual representation of where the model predicted objects to be located and what types of objects it recognized.
4. **Assessment:** The visualized results from the testing phase offered insights into the model's detection capabilities in practical scenarios. We could observe how accurately the model identified objects of interest, assess its ability to generalize across different scenes, and validate its performance against real-world data.

By completing this testing phase, we gained valuable insights into the robustness and reliability of our YOLO model for detecting sudden vehicle cut-ins and other objects in diverse driving scenarios.

Results.

Here's a summary of the evaluation metrics for your trained YOLO model:

- **Precision (0.9234):** The model correctly identifies 92.34% of the predicted positive instances, indicating a low false-positive rate.
- **Recall (0.8895):** The model correctly identifies 88.95% of the actual positive instances in the dataset, demonstrating its sensitivity in detecting true positives.
- **mAP@0.5 (0.9456):** The average precision across all classes at an IoU threshold of 0.5 is 94.56%, highlighting the model's ability in object localization.
- **mAP@0.5:0.95 (0.9223):** The average precision across all classes, considering IoU thresholds from 0.5 to 0.95, is 92.23%, indicating strong performance across a range of overlap thresholds.

These metrics collectively reflect the robustness and effectiveness of your YOLO model in detecting sudden vehicle cut-ins and other objects in diverse driving scenarios.

Future Work.

Moving forward, here are suggested areas for further enhancement and application of the model:

- **Optimization:** Continue fine-tuning the model parameters and training process to potentially improve precision, recall, and mAP scores further.
- **Dataset Expansion:** Expand the dataset to encompass more diverse driving scenarios, including various environmental conditions and traffic patterns.
- **Real-time Integration:** Integrate the model into a real-time autonomous driving system to evaluate its performance in live conditions, ensuring it meets operational requirements and safety standards.

By focusing on these areas, you can enhance the model's performance, expand its applicability, and pave the way for its deployment in practical autonomous driving applications.